

Time_Series_AirPassenger

October 30, 2017

1 Steps to Tackle a Time Series Problem (with Codes in Python)

Note: These are just the codes from article

1.1 Loading and Handling TS in Pandas

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from matplotlib.pyplot import rcParams
rcParams['figure.figsize'] = 15, 6
```

```
In [2]: #Note: aim is not to teach stock price forecasting. It's a very complex domain and I have
data = pd.read_csv('AirPassengers.csv')
print data.head()
print '\n Data Types:'
print data.dtypes
```

	Month	#Passengers
0	1949-01	112
1	1949-02	118
2	1949-03	132
3	1949-04	129
4	1949-05	121

```
Data Types:
Month      object
#Passengers  int64
dtype: object
```

Reading as datetime format:

```
In [3]: dateparse = lambda dates: pd.datetime.strptime(dates, '%Y-%m')
# dateparse('1962-01')
data = pd.read_csv('AirPassengers.csv', parse_dates='Month', index_col='Month', date_parser=dateparse)
print data.head()
```

Month	#Passengers
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121

```
In [4]: #check datatype of index
data.index
```

```
Out[4]: DatetimeIndex(['1949-01-01', '1949-02-01', '1949-03-01', '1949-04-01',
                        '1949-05-01', '1949-06-01', '1949-07-01', '1949-08-01',
                        '1949-09-01', '1949-10-01',
                        ...,
                        '1960-03-01', '1960-04-01', '1960-05-01', '1960-06-01',
                        '1960-07-01', '1960-08-01', '1960-09-01', '1960-10-01',
                        '1960-11-01', '1960-12-01'],
                        dtype='datetime64[ns]', name=u'Month', length=144, freq=None)
```

```
In [5]: #convert to time series:
ts = data['#Passengers']
ts.head(10)
```

```
Out[5]: Month
1949-01-01    112
1949-02-01    118
1949-03-01    132
1949-04-01    129
1949-05-01    121
1949-06-01    135
1949-07-01    148
1949-08-01    148
1949-09-01    136
1949-10-01    119
Name: #Passengers, dtype: int64
```

1.1.1 Indexing TS arrays:

```
In [6]: #1. Specific the index as a string constant:
ts['1949-01-01']
```

```
Out[6]: 112
```

```
In [7]: #2. Import the datetime library and use 'datetime' function:
from datetime import datetime
ts[datetime(1949,1,1)]
```

```
Out[7]: 112
```

2 Get range:

```
In [8]: #1. Specify the entire range:  
ts['1949-01-01':'1949-05-01']
```

```
Out[8]: Month  
1949-01-01    112  
1949-02-01    118  
1949-03-01    132  
1949-04-01    129  
1949-05-01    121  
Name: #Passengers, dtype: int64
```

```
In [9]: #2. Use ':' if one of the indices is at ends:  
ts[:'1949-05-01']
```

```
Out[9]: Month  
1949-01-01    112  
1949-02-01    118  
1949-03-01    132  
1949-04-01    129  
1949-05-01    121  
Name: #Passengers, dtype: int64
```

Note: ends included here

```
In [10]: #All rows of 1962:  
ts['1949']
```


```
Out[10]: Month  
1949-01-01    112  
1949-02-01    118  
1949-03-01    132  
1949-04-01    129  
1949-05-01    121  
1949-06-01    135  
1949-07-01    148  
1949-08-01    148  
1949-09-01    136  
1949-10-01    119  
1949-11-01    104  
1949-12-01    118  
Name: #Passengers, dtype: int64
```

3 Checking for stationarity

3.1 Plot the time-series

```
In [11]: plt.plot(ts)
```

Out[11]: [<matplotlib.lines.Line2D at 0x106405e10>]



3.1.1 Function for testing stationarity


```
In [12]: from statsmodels.tsa.stattools import adfuller
def test_stationarity(timeseries):

    #Determining rolling statistics
    rolmean = pd.rolling_mean(timeseries, window=12)
    rolstd = pd.rolling_std(timeseries, window=12)

    #Plot rolling statistics:
    orig = plt.plot(timeseries, color='blue',label='Original')
    mean = plt.plot(rolmean, color='red', label='Rolling Mean')
    std = plt.plot(rolstd, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show(block=False)

    #Perform Dickey-Fuller test:
    print 'Results of Dickey-Fuller Test:'
    dfctest = adfuller(timeseries, autolag='AIC')
    dfcoutput = pd.Series(dfctest[0:4], index=['Test Statistic','p-value','#Lags Used','Number of Observations'])
    for key,value in dfcoutput[4:].items():
        dfcoutput['Critical Value (%s)'%key] = value
    print dfcoutput

In [13]: test_stationarity(ts)
```



Results of Dickey-Fuller Test:


Test Statistic	0.815369
p-value	0.991880
#Lags Used	13.000000
Number of Observations Used	130.000000
Critical Value (5%)	-2.884042
Critical Value (1%)	-3.481682
Critical Value (10%)	-2.578770
dtype:	float64

4 Making TS Stationary

4.1 Estimating & Eliminating Trend

```
In [14]: ts_log = np.log(ts)
plt.plot(ts_log)
```

```
Out[14]: [<matplotlib.lines.Line2D at 0x10e105210>]
```




4.2 Smoothing:

4.2.1 Moving average

```
In [15]: moving_avg = pd.rolling_mean(ts_log,12)
plt.plot(ts_log)
plt.plot(moving_avg, color='red')
```

```
Out[15]: [<matplotlib.lines.Line2D at 0x10e847b90>]
```



```
In [16]: ts_log_moving_avg_diff = ts_log - moving_avg
ts_log_moving_avg_diff.head(12)
```

```
Out[16]: Month
1949-01-01      NaN
```

```

1949-02-01      NaN
1949-03-01      NaN
1949-04-01      NaN
1949-05-01      NaN
1949-06-01      NaN
1949-07-01      NaN
1949-08-01      NaN
1949-09-01      NaN
1949-10-01      NaN
1949-11-01      NaN
1949-12-01    -0.065494
Name: #Passengers, dtype: float64

```

```

In [17]: ts_log_moving_avg_diff.dropna(inplace=True)
         ts_log_moving_avg_diff.head()

```

```


Out[17]: Month
1949-12-01    -0.065494
1950-01-01    -0.093449
1950-02-01    -0.007566
1950-03-01     0.099416
1950-04-01     0.052142
Name: #Passengers, dtype: float64

```

```

In [18]: test_stationarity(ts_log_moving_avg_diff)

```



Results of Dickey-Fuller Test:

```

Test Statistic      -3.162908
p-value              0.022235
#Lags Used           13.000000
Number of Observations Used  119.000000


```

```
Critical Value (5%)          -2.886151
Critical Value (1%)          -3.486535
Critical Value (10%)         -2.579896
dtype: float64
```


4.2.2 Exponentially Weighted Moving Average

```
In [19]: expwighted_avg = pd.ewma(ts_log, halflife=12)
plt.plot(ts_log)
plt.plot(expwighted_avg, color='red')
# expwighted_avg.plot(style='k--')
```

```
Out[19]: [<matplotlib.lines.Line2D at 0x10ebfa150>]
```



```
In [20]: ts_log_ewma_diff = ts_log - expwighted_avg
test_stationarity(ts_log_ewma_diff)
```





```
Results of Dickey-Fuller Test:
Test Statistic      -3.601262
p-value             0.005737
#Lags Used          13.000000
Number of Observations Used  130.000000
Critical Value (5%)  -2.884042
Critical Value (1%)  -3.481682
Critical Value (10%) -2.578770
dtype: float64
```

4.3 Eliminating Trend and Seasonality


4.3.1 Differencing:

```
In [21]: #Take first difference:
        ts_log_diff = ts_log - ts_log.shift()
        plt.plot(ts_log_diff)

Out[21]: [<matplotlib.lines.Line2D at 0x10ec4f250>]
```



```
In [22]: ts_log_diff.dropna(inplace=True)
        test_stationarity(ts_log_diff)
```



Results of Dickey-Fuller Test:

Test Statistic	-2.717131
p-value	0.071121
#Lags Used	14.000000
Number of Observations Used	128.000000
Critical Value (5%)	-2.884398
Critical Value (1%)	-3.482501
Critical Value (10%)	-2.578960
dtype:	float64

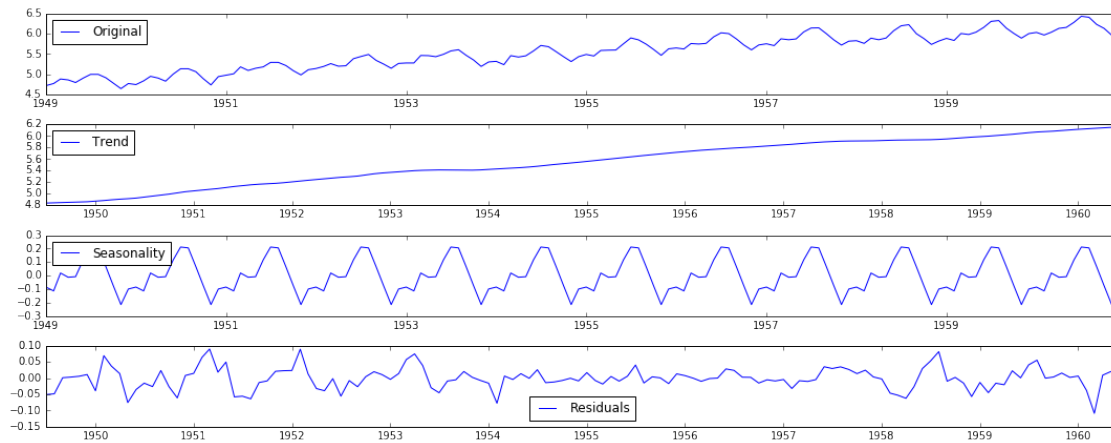
4.3.2 Decomposition:

```
In [23]: from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(ts_log)


trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.subplot(411)
plt.plot(ts_log, label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
```

```
plt.legend(loc='best')
plt.tight_layout()
```



```
In [24]: ts_log_decompose = residual
         ts_log_decompose.dropna(inplace=True)
         test_stationarity(ts_log_decompose)
```



Results of Dickey-Fuller Test:

Test Statistic	-6.332387e+00
p-value	2.885059e-08
#Lags Used	9.000000e+00
Number of Observations Used	1.220000e+02
Critical Value (5%)	-2.885538e+00
Critical Value (1%)	-3.485122e+00
Critical Value (10%)	-2.579569e+00
dtype: float64	

5 Final Forecasting

```
In [25]: from statsmodels.tsa.arima_model import ARIMA
```

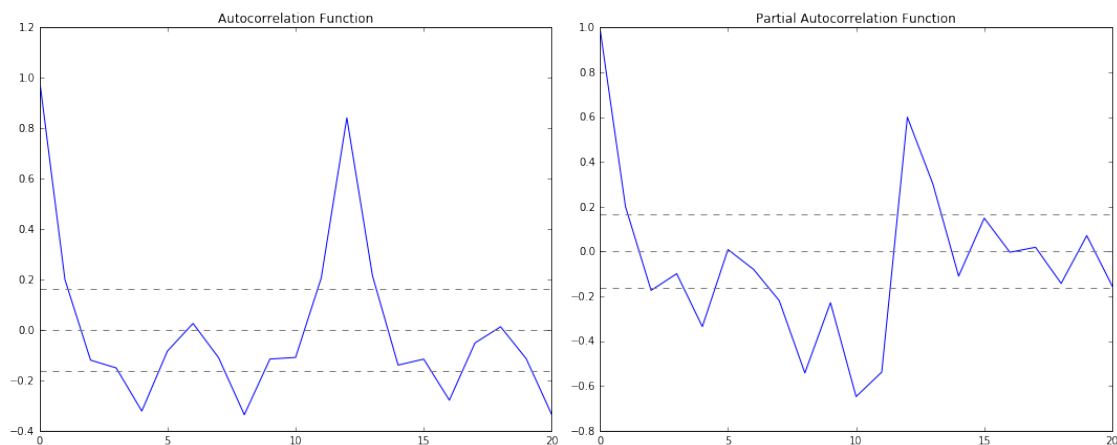
5.0.1 ACF & PACF Plots

```
In [26]: #ACF and PACF plots:
         from statsmodels.tsa.stattools import acf, pacf

         lag_acf = acf(ts_log_diff, nlags=20)
         lag_pacf = pacf(ts_log_diff, nlags=20, method='ols')

         #Plot ACF:
         plt.subplot(121)
         plt.plot(lag_acf)
         plt.axhline(y=0,linestyle='--',color='gray')
         plt.axhline(y=-1.96/np.sqrt(len(ts_log_diff)),linestyle='--',color='gray')
         plt.axhline(y=1.96/np.sqrt(len(ts_log_diff)),linestyle='--',color='gray')
         plt.title('Autocorrelation Function')

         #Plot PACF:
         plt.subplot(122)
         plt.plot(lag_pacf)
         plt.axhline(y=0,linestyle='--',color='gray')
         plt.axhline(y=-1.96/np.sqrt(len(ts_log_diff)),linestyle='--',color='gray')
         plt.axhline(y=1.96/np.sqrt(len(ts_log_diff)),linestyle='--',color='gray')
         plt.title('Partial Autocorrelation Function')
         plt.tight_layout()
```



5.0.2 AR Model:


```
In [27]: #MA model:
         model = ARIMA(ts_log, order=(2, 1, 0))
```

```

results_AR = model.fit(dis=-1)
plt.plot(ts_log_diff)
plt.plot(results_AR.fittedvalues, color='red')
plt.title('RSS: %.4f'% sum((results_AR.fittedvalues-ts_log_diff)**2))

```

Out [27]: <matplotlib.text.Text at 0x1103cae50>




5.03 MA Model

```

In [28]: model = ARIMA(ts_log, order=(0, 1, 2))
results_MA = model.fit(dis=-1)
plt.plot(ts_log_diff)
plt.plot(results_MA.fittedvalues, color='red')
plt.title('RSS: %.4f'% sum((results_MA.fittedvalues-ts_log_diff)**2))

```


Out [28]: <matplotlib.text.Text at 0x1106b3e50>



5.0.4 ARIMA Model:

```
In [29]: model = ARIMA(ts_log, order=(2, 1, 2))
         results_ARIMA = model.fit(disp=-1)
         plt.plot(ts_log_diff)
         plt.plot(results_ARIMA.fittedvalues, color='red')
         plt.title('RSS: %.4f'% sum((results_ARIMA.fittedvalues-ts_log_diff)**2))
```

Out[29]: <matplotlib.text.Text at 0x11077b550>



5.0.5 Convert to original scale:

```
In [30]: predictions_ARIMA_diff = pd.Series(results_ARIMA.fittedvalues, copy=True)
         print predictions_ARIMA_diff.head()
```

```
Month
1949-02-01    0.009580
1949-03-01    0.017491
1949-04-01    0.027670
1949-05-01   -0.004521
1949-06-01   -0.023889
dtype: float64
```

```
In [31]: predictions_ARIMA_diff_cumsum = predictions_ARIMA_diff.cumsum()
         print predictions_ARIMA_diff_cumsum.head()
```

```
Month
1949-02-01    0.009580
1949-03-01    0.027071
1949-04-01    0.054742
1949-05-01    0.050221
```


```
1949-06-01    0.026331
dtype: float64
```

```
In [32]: predictions_ARIMA_log = pd.Series(ts_log.ix[0], index=ts_log.index)
        predictions_ARIMA_log = predictions_ARIMA_log.add(predictions_ARIMA_diff_cumsum, fill_value=0)
        predictions_ARIMA_log.head()
```

```
Out[32]: Month
1949-01-01    4.718499
1949-02-01    4.728079
1949-03-01    4.745570
1949-04-01    4.773241
1949-05-01    4.768720
dtype: float64
```

```
In [33]: plt.plot(ts_log)
        plt.plot(predictions_ARIMA_log)
```

```
Out[33]: [<matplotlib.lines.Line2D at 0x1106756d0>]
```



```
In [34]: predictions_ARIMA = np.exp(predictions_ARIMA_log)
        plt.plot(ts)
        plt.plot(predictions_ARIMA)
        plt.title('RMSE: %.4f'% np.sqrt(sum((predictions_ARIMA-ts)**2)/len(ts)))
```

```
Out[34]: <matplotlib.text.Text at 0x110a6a550>
```

