

CONSISTENCY IN DISTRIBUTED noSQL STORES

Question 1:

The figure shows twelve server machines that together implement a single noSQL store. The servers that are written to and the servers that are read from in the figure share a single node in common, node C. What would happen if node C crashes?

If C crashes, $N_r = 2$, $N_w = 9$ and $N = 11$. But we need $N_r + N_w > N$ to prevent read-write conflicts. So if C crashes and nothing is done this distributed system will no longer be able to prevent read-write conflicts.

To restore the system to a setup that prevents read-write conflicts we need to add at least one node to at least one of the quorums, as demonstrated below.

Let N_{ra} be the number of additional nodes we will add to the read quorum.

Let N_{wa} be the number of additional nodes we will add to the write quorum.

$$N_w + N_{wa} + N_r + N_{ra} > 11$$

So, $N_{wa} + N_{ra} > 0$, which means *at a minimum* we need to choose one of two options:

1. Increase N_r by selecting at least one node from D-L to add to the reading quorum ($N_{ra}=1$, $N_{wa}=0$)
2. Increase N_w by assigning A or B to the write quorum ($N_{ra}=0$, $N_{wa}=1$)

Of course, we could also use higher values for N_{ra} and N_{wa} .

Question 2:

See (b) in figure. Is this a valid quorum configuration? If not, what type of conflict can occur? If so, what are the advantages? What are the drawbacks?

In the example (b) in the figure $N_w = 6$ and $N = 12$. This configuration is not a valid configuration because the rule that $N_w > N/2$ does not hold. This rule is required to prevent write-write conflicts, so in the configuration presented in example (b) write-write conflicts can occur.

Question 3:

See (c) in figure. Is this a valid quorum configuration? If not, what type of conflict can occur? If so, what are the advantages? What are the drawbacks?

In this configuration $N_w = N = 12$ and $N_r = 1$.

This is a valid configuration because (1) $N_r + N_w > N$ and (2) $N_w > N/2$ both hold.

The advantage of having only one node in the read quorum is that this setup enables fast reads; only node F needs to be accessed for read operations. On the other hand, this makes the system vulnerable to failure of a single node.

Regarding the write quorum we have a setup where $N_w = N$. The advantage of having this setup is a high level of consistency; all nodes will have the same data before any write can be confirmed. The drawback is that writes will take a longer time to complete and consequently a high number of write requests may fail.